

wine_testing

October 18, 2022

1 CAPSTONE PROJECT

1.0.1 Summary:

For the wine tasting problem, I have dealt with multicollinearity through PCA and have trained Logistic regression, Random-forest and SVM. I have also run hyperparameter tuning with k-fold cross validation.

```
[1235]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import warnings

warnings.filterwarnings('ignore')
```

```
[1236]: #reading the data

df_entire = pd.read_csv('QualityPrediction.csv')
print(df_entire.shape)
df_entire.head(n=10)
```

(1599, 12)

```
[1236]:   fixed acidity  volatile acidity  citric acid  residual sugar  chlorides \
0           7.4             0.70         0.00           1.9       0.076
1           7.8             0.88         0.00           2.6       0.098
2           7.8             0.76         0.04           2.3       0.092
3          11.2             0.28         0.56           1.9       0.075
4           7.4             0.70         0.00           1.9       0.076
5           7.4             0.66         0.00           1.8       0.075
6           7.9             0.60         0.06           1.6       0.069
7           7.3             0.65         0.00           1.2       0.065
8           7.8             0.58         0.02           2.0       0.073
9           7.5             0.50         0.36           6.1       0.071

   free sulfur dioxide  total sulfur dioxide  density  pH  sulphates \
0              11.0             34.0    0.9978  3.51         0.56
1              25.0             67.0    0.9968  3.20         0.68
```

2	15.0	54.0	0.9970	3.26	0.65
3	17.0	60.0	0.9980	3.16	0.58
4	11.0	34.0	0.9978	3.51	0.56
5	13.0	40.0	0.9978	3.51	0.56
6	15.0	59.0	0.9964	3.30	0.46
7	15.0	21.0	0.9946	3.39	0.47
8	9.0	18.0	0.9968	3.36	0.57
9	17.0	102.0	0.9978	3.35	0.80

	alcohol	quality
0	9.4	5
1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5
5	9.4	5
6	9.4	5
7	10.0	7
8	9.5	7
9	10.5	5

1.1 1. Preprocessing & Exploratory data analysis

1. Treatment of missing values
2. Data format Handling
3. Data visualization
4. Treatment of outliers

1.1.1 1.1 Treatment of missing values

```
[1237]: #percentage of missing values in each column

(df_entire.isnull().sum()/df_entire.isnull().count()).
↳sort_values(ascending=False)
```

```
[1237]: quality          0.0
alcohol                0.0
sulphates              0.0
pH                    0.0
density               0.0
total sulfur dioxide   0.0
free sulfur dioxide    0.0
chlorides              0.0
residual sugar         0.0
citric acid            0.0
volatile acidity       0.0
fixed acidity          0.0
```


dtype: float64

Inference: There are no missing values in the data

1.1.2 1.2 Correlation coefficient measure

```
[1238]: corr= df_entire.corr()  
plt.figure(figsize=(10,6))  
sns.heatmap(corr, cmap='Blues', annot=True)
```

[1238]: <AxesSubplot:>



1.1.3 From the Correlation heatmap it is apparent that features have low-to-medium correlation. We will perform PCA little further before random forest classifier to handle multi-collinearity.

1.1.4 1.3 Data format handling

```
[1298]: df_entire.dtypes
```

```
[1298]: fixed acidity      float64  
volatile acidity      float64
```

citric acid	float64
residual sugar	float64
chlorides	float64
free sulfur dioxide	float64
total sulfur dioxide	float64
density	float64
pH	float64
sulphates	float64
alcohol	float64
quality	int64
dtype:	object

No Data format handling required.


1.1.5 1.4 Data Visualization

```
[1299]: df_entire.columns
```

```
[1299]: Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',
            'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',
            'pH', 'sulphates', 'alcohol', 'quality'],
            dtype='object')
```

```
[1300]: df_entire['quality'].plot.hist()
```

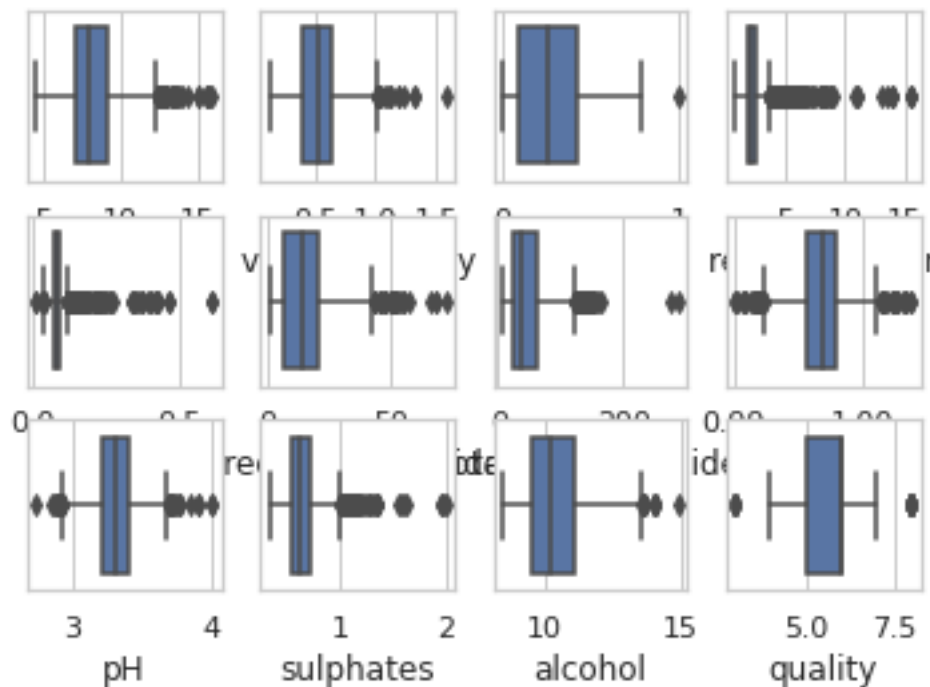
```
[1300]: <AxesSubplot:ylabel='Frequency'>
```



```
[1301]: #sns.boxplot(x="variable", y="value", data=pd.melt(df_entire))
#warnings.filterwarnings('ignore')
```

```
f, axes = plt.subplots(3,4)
sns.boxplot(df_entire['fixed acidity'], ax=axes[0,0])
sns.boxplot(df_entire['volatile acidity'], ax=axes[0,1])
sns.boxplot(df_entire['citric acid'], ax=axes[0,2])
sns.boxplot(df_entire['residual sugar'], ax=axes[0,3])
sns.boxplot(df_entire['chlorides'], ax=axes[1,0])
sns.boxplot(df_entire['free sulfur dioxide'], ax=axes[1,1])
sns.boxplot(df_entire['total sulfur dioxide'], ax=axes[1,2])
sns.boxplot(df_entire['density'], ax=axes[1,3])
sns.boxplot(df_entire['pH'], ax=axes[2,0])
sns.boxplot(df_entire['sulphates'], ax=axes[2,1])
sns.boxplot(df_entire['alcohol'], ax=axes[2,2])
sns.boxplot(df_entire['quality'], ax=axes[2,3])
```

```
[1301]: <AxesSubplot:xlabel='quality'>
```



```
[1302]: df_X = df_entire.drop(columns=['quality'])
df_Y = df_entire[['quality']]
df_X, df_Y
```

```
[1302]: (      fixed acidity  volatile acidity  citric acid  residual sugar  chlorides
\
0          7.4          0.700          0.00          1.9          0.076
1          7.8          0.880          0.00          2.6          0.098
2          7.8          0.760          0.04          2.3          0.092
3         11.2          0.280          0.56          1.9          0.075
4          7.4          0.700          0.00          1.9          0.076
...
1594         6.2          0.600          0.08          2.0          0.090
1595         5.9          0.550          0.10          2.2          0.062
1596         6.3          0.510          0.13          2.3          0.076
1597         5.9          0.645          0.12          2.0          0.075
1598         6.0          0.310          0.47          3.6          0.067
```

```
      free sulfur dioxide  total sulfur dioxide  density  pH  sulphates  \
0          11.0          34.0  0.99780  3.51          0.56
1          25.0          67.0  0.99680  3.20          0.68
2          15.0          54.0  0.99700  3.26          0.65
3          17.0          60.0  0.99800  3.16          0.58
4          11.0          34.0  0.99780  3.51          0.56
...
1594         32.0          44.0  0.99490  3.45          0.58
1595         39.0          51.0  0.99512  3.52          0.76
1596         29.0          40.0  0.99574  3.42          0.75
1597         32.0          44.0  0.99547  3.57          0.71
1598         18.0          42.0  0.99549  3.39          0.66
```

```
      alcohol
0          9.4
1          9.8
2          9.8
3          9.8
4          9.4
...
1594        10.5
1595        11.2
1596        11.0
1597        10.2
1598        11.0
```

```
[1599 rows x 11 columns],
      quality
0          5
1          5
2          5
3          6
4          5
```

```
...      ...
1594      5
1595      6
1596      6
1597      5
1598      6
```

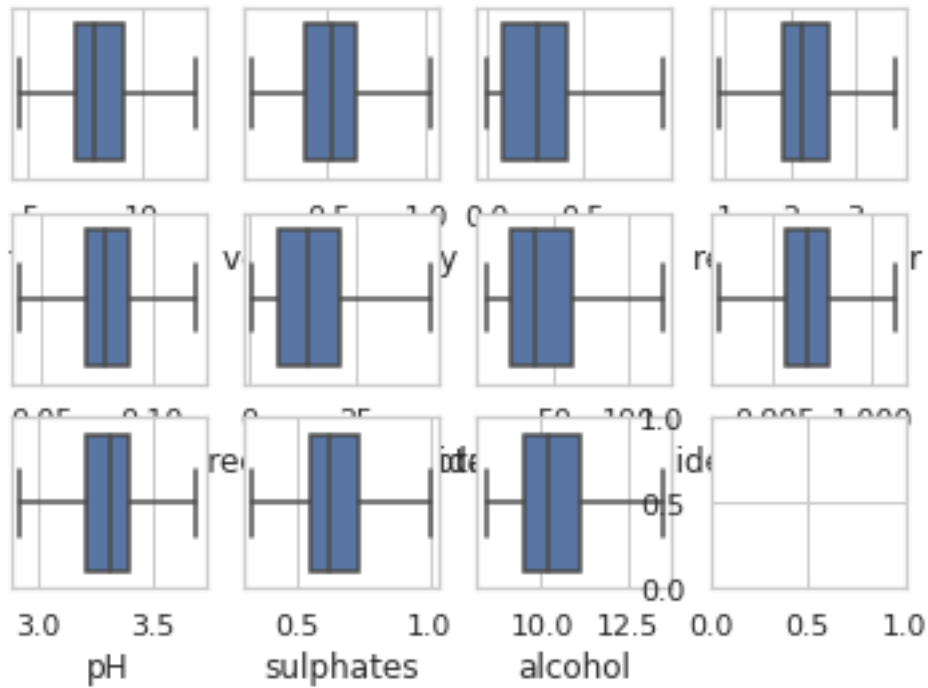
```
[1599 rows x 1 columns])
```

```
[1303]: Q1 = df_X.quantile(0.25)
Q3 = df_X.quantile(0.75)
IQR = Q3 - Q1
lcut = Q1 - 1.5*IQR
rcut = Q3 + 1.5*IQR
L_columns = list(df_X.columns)

for col in L_columns:
    df_X[col].loc[df_X[col] < lcut[col]] = lcut[col] #np.median(df_X[col])
    df_X[col].loc[df_X[col] > rcut[col]] = rcut[col] #np.median(df_X[col])
```

```
[1304]: f, axes = plt.subplots(3,4)
sns.boxplot(df_X['fixed acidity'], ax=axes[0,0])
sns.boxplot(df_X['volatile acidity'], ax=axes[0,1])
sns.boxplot(df_X['citric acid'], ax=axes[0,2])
sns.boxplot(df_X['residual sugar'], ax=axes[0,3])
sns.boxplot(df_X['chlorides'], ax=axes[1,0])
sns.boxplot(df_X['free sulfur dioxide'], ax=axes[1,1])
sns.boxplot(df_X['total sulfur dioxide'], ax=axes[1,2])
sns.boxplot(df_X['density'], ax=axes[1,3])
sns.boxplot(df_X['pH'], ax=axes[2,0])
sns.boxplot(df_X['sulphates'], ax=axes[2,1])
sns.boxplot(df_X['alcohol'], ax=axes[2,2])
```

```
[1304]: <AxesSubplot:xlabel='alcohol'>
```



1.2 2. Modeling

1. Train-Test split
2. Scaling
3. Applying Logistic Regression
4. Applying Random-Forest Classifier
5. Applying SVM Classifier

1.2.1 2.1 Train-Test split

```
[1305]: from sklearn.model_selection import train_test_split
```

```
[1306]: df_X_revised = df_X.drop(columns=['pH', 'residual sugar', 'free sulfur_dioxide']) #, 'fixed acidity', 'citric acid', 'chlorides', 'density', 'total_sulfur dioxide'])
```

```
[1307]: x_train_unscaled, x_test_unscaled, y_train, y_test = train_test_split(df_X_revised, df_Y, test_size=0.4, random_state=111, stratify=df_Y)
```


1.3 2.2 Scaling

```
[1308]: # Robust Scaler  
        ## (x-Q1)/(Q3-Q1)  
        ## Fit for data with outliers  
  
        from sklearn.preprocessing import RobustScaler  
  
        robust_scaler = RobustScaler()  
        x_train = robust_scaler.fit_transform(x_train_unscaled)  
        #df_X = pd.DataFrame(x_train, columns=['fixed acidity', 'volatile acidity',  
        ↪ 'citric acid', 'residual sugar', 'chlorides', 'free sulfur dioxide', 'total_  
        ↪ sulfur dioxide', 'density', 'pH', 'sulphates', 'alcohol'])  
        #df_X.plot.kde()
```

```
[1309]: x_test = robust_scaler.transform(x_test_unscaled)
```


1.3.1 2.2 Applying Logistic Regression

```
[1310]: from sklearn.linear_model import LogisticRegression  
        from sklearn.metrics import confusion_matrix, accuracy_score, roc_auc_score,  
        ↪ plot_confusion_matrix, classification_report
```

```
[1311]: logReg = LogisticRegression(random_state=111)  
        logReg.fit(x_train, y_train)
```

```
[1311]: LogisticRegression(random_state=111)
```

```
[1312]: #prediction and evaluation  
        y_pred = logReg.predict(x_test)  
        plot_confusion_matrix(logReg, x_test, y_test)  
        plt.show()
```



```
[1313]: #accuracy score

print("accuracy score for test data", logReg.score(x_test, y_test))
print("accuracy score for training data", logReg.score(x_train, y_train))
```

```
accuracy score for test data 0.5859375
accuracy score for training data 0.5954118873826904
```

```
[1314]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
3	0.00	0.00	0.00	4
4	0.00	0.00	0.00	21
5	0.63	0.78	0.70	273
6	0.54	0.55	0.55	255
7	0.47	0.29	0.36	80
8	0.00	0.00	0.00	7
accuracy			0.59	640
macro avg	0.27	0.27	0.27	640
weighted avg	0.55	0.59	0.56	640


Inference: Using Logistic regression that the difference between accuracies of training data and test data is about 5% indicating no overfitting

1.4 Performing PCA to handle multicollinearity.

1.4.1 Principal components are used as features to train Random Forest.

```
[1315]: #Applying PCA and apt number of principal components for the dataset.

import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.decomposition import PCA
pca_test = PCA(n_components=8)
pca_test.fit(x_train)
sns.set(style='whitegrid')
plt.plot(np.cumsum(pca_test.explained_variance_ratio_))
plt.xlabel('number of components')
plt.ylabel('cumulative explained variance')
plt.axvline(linewidth=4, color='r', linestyle = '--', x=5, ymin=0, ymax=1)
display(plt.show())
evr = pca_test.explained_variance_ratio_
cvr = np.cumsum(pca_test.explained_variance_ratio_)
pca_df = pd.DataFrame()
pca_df['Cumulative Variance Ratio'] = cvr
pca_df['Explained Variance Ratio'] = evr
display(pca_df.head(10))
```



None

	Cumulative Variance Ratio	Explained Variance Ratio
0	0.331873	0.331873
1	0.557983	0.226110
2	0.685424	0.127441
3	0.802129	0.116706
4	0.888526	0.086397
5	0.956985	0.068459
6	0.981786	0.024801
7	1.000000	0.018214

1.4.2 It can be seen from the PCA that with 5 components about 95% of the variance is covered. We will now use n=5 to get principle components for further training with Random forest classifier.

```
[1316]: pca = PCA(n_components=5)
pca.fit(x_train)

x_train_pca = pca.transform(x_train)
x_test_pca = pca.transform(x_test)
```

1.4.3 2.3 Applying Random Forest Classifier

```
[1317]: from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import mean_squared_error
```

```
[1318]: rf = RandomForestClassifier(random_state=111)
#rf.fit(x_train, y_train)
```

```
[1319]: #hyperparameter tuning
from sklearn.model_selection import RandomizedSearchCV, GridSearchCV

criterion = ['gini']
# Number of trees in random forest
n_estimators = [100]

# Number of features to consider at every split
max_features = [2, 'auto']

# Maximum number of levels in tree
max_depth = range(4,10)
#max_depth.append(None)

# Minimum number of samples required to split a node
min_samples_split = range(2,6)
```

```

# Minimum number of samples required at each leaf node
min_samples_leaf = range(2,5)

# Method of selecting samples for training each tree
bootstrap = [True]

oob_score = [False]

# Create the random grid
random_grid = {'criterion': criterion,
               'n_estimators': n_estimators,
               'max_features': max_features,
               'max_depth': max_depth,
               'min_samples_split': min_samples_split,
               'min_samples_leaf': min_samples_leaf,
               'bootstrap': bootstrap,
               'oob_score': oob_score}
print(random_grid)

```

```

{'criterion': ['gini'], 'n_estimators': [100], 'max_features': [2, 'auto'],
 'max_depth': range(4, 10), 'min_samples_split': range(2, 6), 'min_samples_leaf':
 range(2, 5), 'bootstrap': [True], 'oob_score': [False]}

```

```

[1320]: #rf_random = RandomizedSearchCV(estimator=rf, param_distributions=random_grid,
    ↪n_iter=100, cv=10, n_jobs=-1)
rf_random = GridSearchCV(estimator=rf, param_grid=random_grid, cv=3, n_jobs=-1,
    ↪verbose=2, scoring='accuracy', return_train_score=True)
rf_random.fit(x_train, y_train)

```

Fitting 3 folds for each of 144 candidates, totalling 432 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done 33 tasks      | elapsed: 10.4s
[Parallel(n_jobs=-1)]: Done 154 tasks     | elapsed: 21.1s
[Parallel(n_jobs=-1)]: Done 357 tasks     | elapsed: 41.5s
[Parallel(n_jobs=-1)]: Done 432 out of 432 | elapsed: 49.1s finished

```


```

[1320]: GridSearchCV(cv=3, estimator=RandomForestClassifier(random_state=111),
    n_jobs=-1,
    param_grid={'bootstrap': [True], 'criterion': ['gini'],
               'max_depth': range(4, 10), 'max_features': [2, 'auto'],
               'min_samples_leaf': range(2, 5),
               'min_samples_split': range(2, 6),
               'n_estimators': [100], 'oob_score': [False]},
    return_train_score=True, scoring='accuracy', verbose=2)

```

```
[1321]: sorted_idx = rf_random.best_estimator_.feature_importances_.argsort()
plt.barh(df_X.columns[sorted_idx],rf_random.best_estimator_.
↪feature_importances_[sorted_idx])
```

[1321]: <BarContainer object of 8 artists>



```
[1322]: rf_random.best_params_
```

```
[1322]: {'bootstrap': True,
'criterion': 'gini',
'max_depth': 8,
'max_features': 2,
'min_samples_leaf': 2,
'min_samples_split': 2,
'n_estimators': 100,
'oob_score': False}
```

```
[1327]: #best model
rfc_best = RandomForestClassifier(random_state=11, max_depth=8, max_features=2,
                                min_samples_leaf=2, min_samples_split=2,
↪bootstrap=True)
rfc_best.fit(x_train_pca, y_train)
print("Best RFC accuracy score on test data: ",rfc_best.score(x_test_pca,
↪y_test))
print("Best RFC accuracy score on training data: ",rfc_best.score(x_train_pca,
↪y_train))
print("Best RFC hyperparameters: ", rfc_best.get_params())
```

```

Best RFC accuracy score on test data: 0.603125
Best RFC accuracy score on training data: 0.8300312825860271
Best RFC hyperparameters: {'bootstrap': True, 'ccp_alpha': 0.0, 'class_weight':
None, 'criterion': 'gini', 'max_depth': 8, 'max_features': 2, 'max_leaf_nodes':
None, 'max_samples': None, 'min_impurity_decrease': 0.0, 'min_impurity_split':
None, 'min_samples_leaf': 2, 'min_samples_split': 2, 'min_weight_fraction_leaf':
0.0, 'n_estimators': 100, 'n_jobs': None, 'oob_score': False, 'random_state':
11, 'verbose': 0, 'warm_start': False}

```

```

[1324]: #prediction
#y_pred_rf = rf_random.best_estimator_.predict(x_test)

```

```

[1325]: #base model
rfc = RandomForestClassifier(random_state=11)
rfc.fit(x_train_pca, y_train)
print("default RFC accuracy score on test data: ", rfc.score(x_test_pca, y_test))
print("default RFC accuracy score on training data: ", rfc.score(x_train_pca,
    ↪ y_train))
print("default RFC hyperparameters: ", rfc.get_params())

```

```

default RFC accuracy score on test data: 0.6375
default RFC accuracy score on training data: 1.0
default RFC hyperparameters: {'bootstrap': True, 'ccp_alpha': 0.0,
'class_weight': None, 'criterion': 'gini', 'max_depth': None, 'max_features':
'auto', 'max_leaf_nodes': None, 'max_samples': None, 'min_impurity_decrease':
0.0, 'min_impurity_split': None, 'min_samples_leaf': 1, 'min_samples_split': 2,
'min_weight_fraction_leaf': 0.0, 'n_estimators': 100, 'n_jobs': None,
'oob_score': False, 'random_state': 11, 'verbose': 0, 'warm_start': False}

```

```

[1326]: #print(classification_report(y_test, y_pred_rf))

```

Inference: Using RF classifier with PCA and with hyperparam tuning, the best accuracy obtained on test data is 60%. Notably, there is about 23% difference b/w accuracy score of test and training data indicating high variance.

Also note that with baseline Random Forest model, the testing and training accuracies are respectively 63.75% and 100%. This indicates overfitting or high-variance.

1.4.4 2.4 Applying SVM Classifier

```

[661]: from sklearn.svm import SVC
from sklearn.decomposition import PCA
from sklearn.pipeline import make_pipeline

#pca = PCA(n_components=5, whiten=True, random_state=42)
svc = SVC(kernel='rbf', class_weight='balanced')
#model = make_pipeline(pca, svc)

```

```
model = make_pipeline(svc)
```

```
[662]: from sklearn.model_selection import GridSearchCV
param_grid = {'svc__C': [1, 5, 10, 50],
              'svc__gamma': [0.0001, 0.0005, 0.001, 0.005, 1.0]}
grid = GridSearchCV(model, param_grid)

%timeit grid.fit(x_train, y_train)
print(grid.best_params_)
```

8.37 s ± 345 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)

{'svc__C': 10, 'svc__gamma': 1.0}

```
[663]: #prediction
model = grid.best_estimator_
y_pred_svc = model.predict(x_test)
```

```
[664]: print("accuracy score for test data: ", model.score(x_test, y_test))
print("accuracy score for training data: ", model.score(x_train, y_train))
```

accuracy score for test data: 0.6125

accuracy score for training data: 0.9601250977326036

```
[665]: print("classification report: ", classification_report(y_test, y_pred_svc))
```

classification report:		precision	recall	f1-score	support
	3	0.00	0.00	0.00	2
	4	0.09	0.09	0.09	11
	5	0.69	0.71	0.70	136
	6	0.60	0.59	0.60	128
	7	0.52	0.55	0.54	40
	8	0.00	0.00	0.00	3
accuracy			0.61		320
macro avg	0.32	0.32	0.32		320
weighted avg	0.60	0.61	0.61		320

Inference: Using SVM classifier, the best accuracy obtained on the test data is about 57%. Notably, there is about 33% difference between the accuracy of test and training data indicating higher variance.

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```